Independent Reviewer Verdict

Scored fresh against the NIH CRP criteria, from the delivered PDF, as a study section would meet it. No credit given for intent, history, or work not on the page.

QUESTION 1

Criterion scores

NIH anchors are absolute, not relative: 1 to 3 is high merit with negligible weaknesses, 4 to 6 moderate, 7 to 9 low. The bar is shown where a skeptical but fair reviewer would land, and the band shows the spread I would expect across three reviewers.

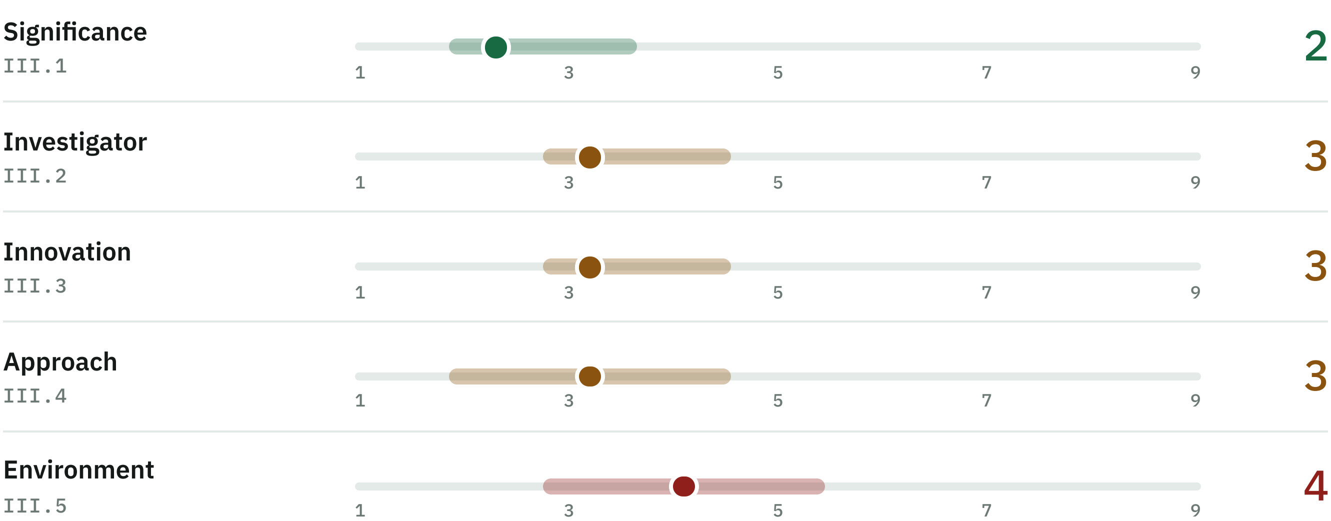

Bar = my score · Band = expected spread across a three-reviewer panel

Significance · 2 **The strongest criterion, and it earns it**

Every subject II.1 requires is present and evidenced: unmet need, market segments and customers, competitive environment, relation to academic and industry efforts, adoption hurdles. The urgency is documented rather than asserted, with $80,080 annual home care cost, $58 billion unclaimed, 63.3 percent of providers declining cases they could not staff, and unmet ADL needs predicting emergency use, readmission, placement, and mortality.

What makes it a 2 rather than a 3 is market-pull demonstrated with the customer's own money. Table 1 shows a 30-caregiver agency spending $11,700 a year to hire against $3,300 for this product. That is the rarest thing in an SBIR Significance section: evidence the customer is already buying the problem's solution at a worse price.

What a skeptic has to overlook to agree. Table 2 gives Olera a checkmark in all five columns while every competitor carries at least one cross. Reviewers are trained to distrust self-comparison tables the applicant wins outright, and two of those checkmarks describe capabilities the Approach says do not exist yet:

execution to established care is what Aim 1 builds, and the staffing product has placed roughly twenty people by hand. A reviewer who catches that reads the table as an assertion of current-state parity that the rest of the document contradicts, and scores 3.

`Investigator · 3` **Credible, but the criterion asks a question the document does not answer**

The oversight mechanics are genuinely well handled, which is unusual. The criterion asks about "robust oversight of CROs and a clear means to integrate their effort," and the document answers directly: Clemson through Dr. Fan at 25 percent effort holding every protocol, integrated through quarterly protocol review and data-management procedures in the subaward; ADC under subcontract with a defined deliverable; the PI holding integration authority and every gate decision against a quarterly milestone review.

The gap. The criterion also asks, in NIH's words, "how successful have the PD(s)/PI(s) been in commercializing technologies and discoveries in the past?" The document's honest answer is that they have not yet. It is also silent on who runs commercial operations. The central risk in this application is operating ten markets in thirty-six months, and no member of the team is identified as owning that.

`Innovation · 3` **One genuinely strong innovation, one reframing, one that reads as automation**

Key Innovation 3 is the best intellectual content in the document. The three-layer argument is precise, and the disqualifying contrast is memorable and correct: "NHATS is a survey; claims data record what was billed, not what was needed and refused." It establishes a real asset with a real moat, and the flywheel argument is stated as a mechanism rather than a slogan.

Key Innovation 1 is a legitimate and non-obvious reframing. That every existing channel competes for the same constrained pool, so one provider's hire is another's vacancy, is true and most readers will not have thought of it. The portable verified experience record is a plausible technical mechanism.

Key Innovation 2 is the weak one. Agents that file applications and follow up is workflow automation, and the criterion asks specifically whether the product is "substantially innovative compared to all existing approaches *as well as those in development*." The document does not engage what is in development elsewhere in agentic workflow automation at all, which in 2026 is a crowded and well-capitalized field. That omission is the clearest thing holding this criterion below a 2.

A bracketed instruction to a team member sits inside Key Innovation 1: `[TJ: confirm the name and scope of the verified worker record.]` A reviewer meeting that inside the flagship innovation reads it as an unfinished application, and it costs more than its size.

`Approach · 3` **Unusually rigorous analytics, with two specific holes and one unevidenced assumption**

What is genuinely strong. Randomization inference as the primary comparison for an eight-cluster design is the correct method and is rarely done correctly at this scale. The blinded expert audit specifies a stratified sample of 120 cases sized to a six-point half-width, Cohen's kappa, inter-rater reliability established before any system output is adjudicated, and audited households held out of tuning data. Aim 2 carries a real power calculation. Retention uses discrete-time survival with competing events and restricted mean survival time feeding lifetime value with a confidence interval. Every aim ends with potential problems, alternative strategies, and a gate, which is exactly what the criterion asks for and is better executed here than in most applications at this program.

Hole one: the primary Aim 3 comparison has no stated precision. The document sizes account-level conversion to a 95 percent half-width near 10 points, then says a biostatistician "confirms that sizing against the conversion rate measured in Aim 2 rather than an assumed one." But the primary comparison is at the market level across four arms with two markets each, and the document never states what interval that comparison delivers. A methodologist will ask whether four price arms can be distinguished at all, and the answer is not on the page.

Hole two: the headline Aim 2 outcome has no performance threshold. Table 5 lists "Households reaching established aid or care" with the criterion "Estimated, ±5 pts." That is a precision target, not a success target. The document is candid about why, arguing no benchmark exists and that inventing one would be worse. Half of reviewers will read that as intellectual honesty; the other half will observe that the central outcome of the middle aim cannot fail.

The unevidenced assumption. Task 3.1 states that "the binding constraint on this aim is operating attention, not capital," and the document then says nothing further about operating attention. No headcount, no per-market staffing model, no evidence that a single market has been run end to end. This is the largest unmitigated risk in the application.

Smaller. Table 4 carries four criteria for a year-long engineering aim, only two of which gate, and Tasks 1.1 through 1.3 have no quantitative criterion of their own, which a solicitation-literal reviewer will note against the requirement for milestones covering each objective. The 85 percent agreement bar is asserted without anchoring to literature or prior performance.

Environment · 4 **Real partners, and the sharpest edge on the application**

The positives are real: a Clemson subaward with defined effort and protocol ownership, ADC and a named LCSW panel lead with letters signposted, and a regulatory position that reduces the regulatory-resource question close to zero. Most valuable of all, two named financial institutions, Ziegler and Equitage Ventures, have done diligence, and the document says that diligence set the milestone tables. That is genuine third-party validation and a reviewer will credit it.

The edge. The criterion asks directly, "how successful is the company's track record in commercializing prior SBIR/STTR projects?" Nine years, two NIA awards, and the Progress Report's own words: "The provider tools launched nationwide in July 2026 with no paywall, so there is no pricing history to read." The total commercial history visible anywhere in the document is three providers paying between $50 and $275. A reviewer inclined to be generous reads free-by-design strategy. A reviewer reading the criterion literally reads a company that has never sold anything applying to a program whose purpose is to bridge companies that are close to selling.

Diligence is also not commitment. The Fundraising criterion tests the ability to secure third-party funds equal to or exceeding the NIH request, which is $4 million, and nothing in this document suggests that is close.

QUESTION 2

Overall impact

> REALISTIC RANGE FROM A STRONG BUT SKEPTICAL STUDY SECTION
>
> ## 24 – 34
>
> **Most likely landing after discussion: 28 to 30.** Comfortably above triage, discussed, and in the middle of the discussed range rather than near the top of it.

Overall impact for this program is defined as the likelihood of "a sustained, powerful influence on **the relevant market offering**." Three things pull against that phrase, and they are not the same as the criterion weaknesses.

First, scale, and it comes from a number the document volunteers. Significance states that at modeled steady state a market carries about 50 paying accounts and $165,000 in annual recurring revenue. A reviewer will multiply. Ten markets is roughly $1.65 million of annual recurring revenue, against a $4 million federal award. The document has an answer, that the markets are an instrument rather than the deliverable and the real scale lives with the institutional customer class, but the instrument argument cuts against impact rather than for it, and the institutional class is explicitly not tested here. Putting the per-market figure in Significance was a decision that helps the value proposition and hurts the impact score, and reviewers will notice both.

Second, execution risk. Ten market launches, two IRB studies, an engineering build, a four-arm pricing experiment, and an independent CPA rebuild, in thirty-six months, by a company whose demonstrated operational record is twenty caregiver placements and three paying providers. Nothing in the document closes the distance between those two facts.

Third, finish. Four visible placeholders, one broken sentence on page 1 where a comma stands in for a period, and three of twenty-six references that are not verifiable as written, including one carrying `[CONFIRM EDITION AND YEAR]` in the bibliography and one systematic review with no authors or journal. Individually trivial. Together they read as an application submitted before it was finished, and reviewers discount for that more than the substance warrants, because it is the one signal they can verify without expertise.

QUESTION 3

How each archetype reads it

The top-to-bottom scientific reader

25

IMPACT

Experiences it as unusually well-written for this program. The three gates, the vicious cycle, and the line that the money to support older adults at home largely exists today but is spent later on worse outcomes, all survive to the summary statement. Figure 1 does real work; Figure 3 makes an abstract product concrete in a way few applications manage.

Trips on the punctuation error early, which is the worst possible place for it, and on the bracketed notes to a colleague. Comes away favorable but noting that the application "reads like a late draft."

The methodologist

30

IMPACT

The most engaged reader in the room, and mostly favorably. Recognizes randomization inference as the right choice for eight clusters, notes the competing-risks framing and RMST, approves of inter-rater reliability established before adjudication and audited households held out of tuning data, and approves of acquisition cost including non-converters.

Then writes three lines that dominate the critique: the market-level primary comparison has no stated precision, the Aim 2 primary outcome has no threshold, and the 85 percent agreement bar is unanchored. On the engineering side, reads the Aim 1 architecture as competent but generic, and reads `[TJ: confirm runtime and model choices.]` as undoing the paragraph it terminates.

The commercialization reviewer

35

IMPACT

The hardest read, and likely the discussant. Genuinely likes Table 1, likes that price is measured rather than assumed, and likes the independent CPA rebuild with discrepancies reported rather than reconciled, which is a credibility move most applicants do not make.

Then: pre-revenue after nine years and roughly $2.7 million of federal funding; $165,000 per market; no named commercial or operations leadership; and the customer class where the money actually is, explicitly untested. Asks the question that sets the tone of the discussion, which is whether this is a commercialization award or a fourth research award.

The solicitation-literal reviewer

28

IMPACT

Goes down II.1 through II.4 and finds every required subject present, which is not a given. II.4 in particular is well served: Table 7, mapping each retired risk to the funding that retired it and naming the one that remains, is a better Progress Report artifact than most applications produce.

Flags that quantitative milestones do not cover each objective in Aim 1, that the Project Management Plan is scored under Approach but lives in another document, that the inclusion statement is present but minimal, and itemizes the placeholders and unverifiable references as completeness defects.

QUESTION 4

Strongest score-driving features

What a supportive reviewer would point to when defending a favorable score.

1. **Table 1.** The single best artifact in the document. It shows the customer already spends between 1.8 and 12 times more on the alternative, on the customer's own numbers. This is what a compelling value proposition looks like on paper.
2. **The structural claim.** That navigation and capacity are jointly necessary and separately insufficient is non-obvious, correct, and it is what makes the workforce product read as necessity rather than scope creep. It is the spine of the application and it holds.
3. **Key Innovation 3.** The three-layer argument, the NHATS and claims contrast, and a moat that compounds with use. The most memorable and most defensible innovation claim in the document.
4. **The gate, alternative, and stop-rule architecture.** Every aim ends with problems, alternatives, and a gate; the month 30 stop rule defines disconfirmation in advance. "A commercialization program should fund projects willing to find out they are wrong" is the kind of sentence a supportive reviewer quotes aloud.
5. **Table 7.** Risks retired, the funding that retired each, and the risk that remains, ending on an unretired row. Precisely what II.4 asks for, executed better than most.
6. **Independent financial validation.** An outside CPA rebuilding the numbers from raw records without access to the operating model, with discrepancies reported rather than reconciled. Few applicants volunteer to be checked.
7. **Measurement discipline throughout.** Acquisition cost includes non-converters; an outcome counts once; profitability flagged descriptive below twenty accounts; the audit sample held out of tuning data. A reviewer reads these as a team that has thought about how it could fool itself.
8. **Visual and prose craft.** Restrained figures, consistent vocabulary, no padding. The document is materially easier to review than its peers, and that earns goodwill.

QUESTION 5

Most consequential attack surfaces

Ranked by how much damage each would do if raised at study section.

1. **Nothing demonstrates the company can operate at the proposed scale.** Ten markets against a record of twenty placements and three paying providers, with the constraint named as "operating attention" and operating attention never addressed. This is the weakness most likely to cap the score, and it is the one a reviewer can raise without any domain expertise.
2. **The primary Aim 3 comparison's precision is unstated.** Four price arms, two markets each, and no statement of what the market-level analysis can resolve. The most technically damaging line available to a methodologist.
3. **Table 2's all-checkmark row asserts capabilities the Approach says do not exist yet.** Specific, verifiable within the document, and it invites a reviewer to question other claims.

4 **Aim 2's headline outcome carries a precision target, not a performance target.** Reads to half the room as honesty and to the other half as an outcome that cannot fail.

5 **Pre-revenue after nine years,** against two scored criteria that ask about commercialization track record in so many words.

6 **The scale of the market offering,** now computable from the document's own figures at roughly $1.65 million of annual recurring revenue across the ten-market footprint.

7 **Aim 1's quantitative milestones do not cover Tasks 1.1 through 1.3,** and the 85 percent agreement bar is unanchored.

8 **Innovation does not engage what is in development elsewhere,** which the criterion asks for explicitly, in a field where a great deal is in development.

9 **Four placeholders, one broken sentence, three unverifiable references.** Cheap to fix, disproportionately expensive to leave.

QUESTION 6

Is it competitive at 10 to 20 as written?

No. It is competitive at 24 to 34 and would most likely land at 28 to 30. Four specific things stand between it and the 10 to 20 range, and only three of them are editorial.

Blocker 1 **The operational scale-up answer does not exist on the page**

A reviewer cannot defend a 15 for an award whose central activity is opening ten markets when the application never says who opens them, how many people it takes, or whether the applicant has ever opened one. The document names the constraint and then changes the subject. This is fixable and it is the highest-yield fix available.

Blocker 2 **The revenue the award produces is small against the ask**

This is the one that is not editorial. At $275 a month, the entire home-care subscription market at full national penetration is a modest business, and the document's own per-market figure makes that computable. The strategic answer is the institutional customer class, and the application deliberately does not test it. That is a defensible research decision and a costly impact decision, and reviewers will feel the tension.

Blocker 3 **Two methodological gaps a strong reviewer will find**

The unstated market-level precision in Aim 3 and the absent performance threshold in Aim 2. Both are answerable in a paragraph each. Left as they are, they hand the methodologist the two sentences that shape the written critique.

Blocker 4 **The document is not finished, and it shows**

Four bracketed notes, a comma where a period belongs on page 1, and three references that would not survive a check. None of this is substantive. All of it is visible, and it licenses a skeptic to assume the parts they cannot check are equally unfinished.

QUESTION 7

What they would say in the room

The favorable case

> *"This is the most carefully constructed application in my pile. The problem is real and urgent and they document it rather than assert it. The insight that navigation without capacity just delivers families to a supply side that cannot absorb them is correct and I had not thought about it that way. And they have found something genuinely defensible in the data asset. Nobody has the record of what actually happened between a family recognizing a need and care starting, and if they build it, they own it."*

> *"What persuades me is the discipline. They put a gate on every aim, they pre-register the pricing analysis, they hire an outside accountant to check their own numbers and commit to reporting discrepancies rather than reconciling them, and they write down in advance the conditions under which they would declare themselves wrong. I do not see that. I would rather fund a team that might fail honestly than one that cannot."*

> *"And the value proposition is the customer's own arithmetic. A thirty-caregiver agency spends eleven thousand seven hundred dollars a year hiring and they are charging three thousand three hundred. That is not a category they have to create."*

The skeptical case

> *"Nine years and two NIA awards and there is no revenue. The Progress Report says the provider tools launched with no paywall so there is no pricing history. That is a strategy, fine, but the CRP is supposed to bridge companies that are nearly selling, and this company has three customers who paid something between fifty and two hundred seventy-five dollars."*

> *"My real problem is operational. They propose ten markets in three years. They have placed twenty caregivers. They say themselves that the binding constraint is operating attention and then they never tell us whose attention, how many people, or what a market launch costs in labor. I do not know who is doing this work."*

> *"And do the arithmetic they give us. A hundred sixty-five thousand a market. Ten markets. That is under two million in recurring revenue for a four million dollar award. They tell us the markets are an instrument, not the deliverable, and I believe them, but then I am funding an instrument, and the impact criterion asks about a powerful influence on the market offering."*

> *"On the design: the primary Aim 3 comparison is four price arms across eight markets and they never say what that comparison can actually resolve. And their Aim 2 headline outcome has a precision criterion, not a performance criterion. I am not sure that aim can fail."*

> *"Also, there are notes to a team member still in the text and a citation that says confirm edition and year. That tells me something."*

QUESTION 8

Funding verdict

This is a good application that is not yet a fundable-this-round application, and the distance between those two things is smaller than it usually is.

It will be discussed. It is well above the line where applications are triaged unread, and the assigned reviewers will engage with it seriously because it is clear, specific, and obviously the work of people who know their field. On the merits I would expect it to land near the middle of the discussed range, around 28 to 30, which in most cycles is short of the payline and close enough that a resubmission is the obvious move rather than a long shot.

What it is genuinely strong at: a documented and urgent problem, a value proposition backed by the customer's existing spend, one innovation with a real and compounding moat, an analytic plan more rigorous than the program typically sees, and a governance structure with gates, alternatives, and a prespecified condition for stopping. Those are not small things and I do not want to bury them under the critique. A reviewer defending this application has real material to work with.

What holds it back is concentrated rather than diffuse. One structural problem, that the company has never demonstrated it can operate at the scale it proposes, and that problem sits directly underneath the aim that carries the commercial argument. One strategic problem, that the revenue this produces is modest relative to the ask, which is a decision about what to test rather than a writing flaw. Two methodological gaps, each answerable in a paragraph. And a finish problem that costs more than it should.

My verdict as a fresh reviewer: score it in the high twenties, argue for it in discussion because the design discipline is real and the problem matters, and say plainly in my critique that I would fund the resubmission if they can tell me who runs ten markets and what that costs in people. That is the sentence I would want the applicant to read.

Scored against NIH CRP review criteria III.1 through III.5 and Research Strategy requirements II.1 through II.4, from the delivered 12-page PDF. Commercialization Plan, Project Management Plan, Letters of Support, and PHS Human Subjects materials assumed present and strong where signposted, and not scored here.