

Mock Summary Statement

Independent NIH CRP Peer Review

Olera, Inc. · SBIR Commercialization Readiness Pilot · Research Strategy and Commercialization Plan · 30 August 2026

(Privileged Communication — Mock Review, Not an NIH Document)

Applicant Organization	Olera, Inc.
Principal Investigator	Falohun, Tokunbo (TJ)
Project Title	CareNavigator + Caregiver Staffing: Executing and Verifying the Care-Establishment Pathway
Mechanism	SBIR Commercialization Readiness Pilot (CRP)
Requested Support	~\$4.0M direct + indirect, 3 years (per Commercialization Plan §6)
Materials Reviewed	Research Strategy (Significance / Innovation / Approach / CRP Progress Report); Commercialization Plan (§§1–9)
Materials Not Reviewed	Specific Aims; Human Subjects & Clinical Trials forms; Letters of Support; biosketches; budget justification; ancillary forms
Review Basis	Per instruction, absent components are assumed competent, internally consistent, and reasonably supportive of what the supplied documents signpost. No criterion is scored down for their absence.

Reviewer note on package scope. The one supplied exemplar summary statement (2R44AG074116-04) was used solely to calibrate genre, tone, structure, and specificity. Its scores, findings, and human-subjects determination were quarantined and exerted no influence on this review. The separately uploaded house-style CRP Progress Report is textually identical to the Progress Report inside the Research Strategy and was treated as one document.

RESUME AND SUMMARY OF DISCUSSION

This CRP application proposes to extend a mature, NIA-funded eldercare navigation platform (CareNavigator) from *planning* into *execution*, and to pair it with a provider-facing workforce product (Caregiver Staffing) that supplies new caregivers to the local labor market. Three aims are proposed in a strictly gated sequence: Aim 1 engineers and technically verifies a longitudinal Care Establishment Model, bounded AI agent execution, a field-learning knowledge layer, and staffing infrastructure, and validates usability and trust with 75 stakeholders before any real-world action; Aim 2 moves the verified automation envelope into live markets and measures whether an index family episode reaches verified aid or care establishment within 90 days; Aim 3 introduces price and tests whether demonstrated value converts into paid conversion, retention, unit economics, and a repeatable market-entry playbook, and assembles an institutional-buyer evidence package.

Enthusiasm centered on four points that reviewers regarded as genuine and well argued. First, the Significance section is among the clearest customer-definition sections the panel would see in this round: it separates user (family, free), participant (provider, free), and buyer (provider for staffing; risk-bearing institutions, evidence-gated) without blurring

them, and it is explicit that the second market is gated on evidence the applicant does not yet have. Second, the central intellectual claim — that the eldercare care-establishment pathway must be made *computable* before it can be reliably automated or measured, and that executed cases produce local operational truth unavailable on the public web — is a substantive and non-obvious thesis rather than an assertion that AI is innovative. Third, the risk-retirement narrative in the Progress Report is cumulative and unusually disciplined: each prior stage tested one uncertainty before the next layer was built, and the applicant is candid that manual delivery of Caregiver Staffing was deliberate learning rather than product. Fourth, the gating architecture — technical verification before human exposure, human-protection overrides, real-world value before price, and an explicit willingness to publish a negative result and stop rather than scale a failing model — is exactly the posture a CRP should exhibit, and several reviewers described it as the strongest structural feature of the application.

Concerns were substantial and concentrated in Approach and in the arithmetic linking the two documents. The most serious is that the application's gates are structurally excellent but numerically empty at precisely the decision points that matter: Aim 1's Care Establishment Model criterion is "prespecified completeness" with no threshold; Aim 2's primary criterion is a day-90 establishment proportion "estimated with prespecified precision" against a success criterion that is never stated, and an N of ~400 that is explicitly deferred to a power analysis not present in the application. A GO/NO-GO gate whose threshold will be set after the application is reviewed is not a gate that reviewers can score. Second, Aim 2's primary endpoint is agnostic to whether the innovation performed the work: a family confirming that care was established is fully consistent with Olera's two existing full-time call-center staff and human navigators executing every consequential step, and no endpoint, criterion, or prespecified target separates agent-executed from human-executed pathways. Third, the workforce revenue model rests on a throughput assumption ("10 successful hires per county per month," described as conservative) that is roughly twenty times the company's entire lifetime placement history and is supported by no data in either document; at the applicant's own reported pilot yield the Year 3 revenue figure requires on the order of seventeen thousand applicants in six months. Fourth, the two documents disagree materially on the number of markets — the Research Strategy's Progress Report says two, the Commercialization Plan says approximately eight and builds its entire revenue model on eight, and Aim 2 says the number will be fixed later — and because market count, sample size, revenue, and operating cost are all functions of one another, this single discrepancy destabilizes the feasibility case in both directions. Fifth, Aim 2 is single-arm with no concurrent comparator, so the institutional value proposition on which the Commercialization Plan's Year 4 and Year 5 revenue depends cannot be causally supported by the evidence the Research Strategy is designed to produce; Aim 3, Task 3.4 in fact delivers a *design* for a future payer study rather than payer-grade evidence, while the Commercialization Plan books a \$250K institutional contract twelve months after the award ends. Reviewers also noted a weak formal IP position (trade secret and copyright only; patents merely "will evaluate"), a Fundraising Plan resting on relationships the applicant itself twice describes as not being commitments, and a very lean operating team — one full-time engineer and two full-time call-center staff — for a program that must build a longitudinal state model, agentic execution with browser automation and fax, a field-learning layer, and staffing infrastructure in Year 1 alone.

Disagreement was expected on two axes. Some reviewers would score Significance a 3 rather than a 2, on the grounds that the SBIR Significance criterion explicitly folds in whether the Commercialization Plan demonstrates a *high probability* of commercialization, and that probability is undercut by unvalidated pricing, a nil commercialization track record across ~\$5.7M of prior federal support, and no private capital history. Others would hold Significance at 2 and locate those concerns entirely in Environment and in the Commercialization Plan assessment. A second disagreement concerns whether the unspecified Aim 2 threshold is a fixable drafting defect or a substantive weakness; the majority view in this simulated panel was that because the same document specifies thresholds precisely elsewhere (SUS ≥ 72 , TIAS ≥ 5 , $\geq 90\%$ action agreement, $\geq 50\%$ placement), the absence of a threshold at the single most consequential gate reads as unresolved rather than merely unwritten.

On balance, the panel judged that the significance of the problem, the coherence of the technical thesis, the discipline of the risk sequencing, and the team's demonstrated delivery against two prior NIA awards outweigh the identified weaknesses, but not decisively. The weaknesses are specific, quantitative, and concentrated in the Approach and in the commercial arithmetic, and several are correctable. Impact is likely to be very good.

CRITIQUE 1 — INTEGRATED PRIMARY REVIEWER

Significance: 2 · Investigator(s): 3 · Innovation: 3 · Approach: 5 · Environment: 4

Overall Impact

This application proposes the right project for the CRP mechanism: it is not further discovery, and it is not a request to fund scale-up of an already-proven business. It sits precisely in the gap the mechanism exists to bridge — a nationally deployed, peer-reviewed, demonstrably used platform that has never been shown to *complete* the job it starts, run by a company with essentially no commercial revenue after ~\$5.7M of federal support. The applicant's framing of that gap is honest and the three-aim sequence maps onto it cleanly: build and verify, then prove real-world value, then test whether value converts to money. The intellectual core — that eldercare cannot be automated until the pathway is represented as computable longitudinal state, and that executing cases generates local operational truth that no directory or claims file contains — is a real contribution and the strongest reason to fund this work.

The application's principal limitation is that it is more rigorous about *architecture* than about *numbers*. The gating logic is superb; the thresholds behind two of the three gates are absent or unjustified. The instrumentation plan is comprehensive; the primary endpoint does not distinguish the intervention from the incumbent human process it is meant to replace. The revenue model is arithmetically self-consistent and unusually candid in labeling hypotheses as hypotheses; the single throughput input that drives it is ~20× the company's demonstrated performance and is called conservative. The Commercialization Plan needs causal evidence of avoided utilization to justify its institutional revenue; the Research Strategy explicitly and correctly declines to produce it. These are not fatal, but they are the difference between an application a panel funds enthusiastically and one it funds with reservations.

Score-driving contributors, positive: an exceptionally well-specified unmet need with a clean user/participant/buyer separation; a genuine and defensible technical thesis in the Care Establishment Model; a disciplined, honest, cumulative risk-retirement history including real (if tiny) revenue and 900 workforce applicants; and a gated design with human protection as an explicit override. Score-driving contributors, negative: unspecified success criteria at the Aim 2 gate and no power justification for $N \approx 400$; a primary endpoint blind to automation share; a workforce throughput assumption unsupported by and contradicted by the applicant's own pilot yield; a market count that differs materially between the two documents; a single-arm design that cannot support the institutional value claim the Commercialization Plan monetizes; and a very lean team relative to the Year 1 engineering load.

1. Significance — Score: 2

Strengths - The unmet need is established with specificity and appropriate citation: \$80,080 annual full-time home care cost, Medicare's exclusion of custodial care, ~\$58B in unclaimed aid, 63.3% of home-care providers declining cases for lack of staff in 2023, and the demographic scissors of 82M over-65s by 2050 against family caregivers per adult over 80 falling from >7 to 4. - The link from unmet need to health consequence is made with real evidence rather than assertion: unmet ADL needs independently predicting ED use, readmission, nursing-home placement, and mortality, and the

Commercialization Plan's HCBS figures (52% vs. 34% ED use; 36% vs. 24% hospital/rehab stays). - The customer architecture is unusually clear and is consistent across both documents. Families are free users; providers are free network participants *and* paying staffing customers; risk-bearing institutions are explicitly labeled an evidence-gated market. Very few applications of this type keep user, participant, and buyer this cleanly separated. - The applicant states plainly that the institutional market is *not yet purchasable* and that the CRP must generate the evidence to open it. This is the correct posture and reviewers credited it. - The competitive analysis (Research Strategy Table 2) is honest and structurally intelligent: alternatives are mapped against the seven steps of the care-establishment pathway rather than against a vague category, and Olera is given full marks only in the "after the CRP" row. Named competitors include the ones that actually matter (Wellthy, Homethrive, Cariloop, A Place for Mom, Caring.com, CareYaya, Indeed, and general-purpose AI). - Hurdles to adoption are named for each of the four constituencies rather than generically, and each is mapped to a CRP activity (Commercialization Plan Table 5). - The GUIDE Model is correctly deployed: not as proof that Olera will be paid, but as precedent that CMS is already paying for navigation and caregiver support with the explicit aim of delaying institutionalization.

Weaknesses - The beachhead market is never sized in dollars. The application gives worker counts (4.68M aides), openings (760,500/yr), and turnover (75%), but never the addressable *spend* on caregiver recruitment that Caregiver Staffing competes for. Without it, reviewers cannot judge whether a \$250-per-hire product addresses a large market or a thin one. - The application's own scale illustration argues against it. At 500 counties, 60,000 successful hires — approximately 8% of all annual U.S. direct-care openings — yields only \$15M in revenue. Capturing 8% of a national labor market for \$15M suggests the price point, the value capture, or the market framing is mismatched, and reviewers will notice that the illustration meant to demonstrate headroom instead demonstrates a low revenue ceiling per unit of operational effort. - The strategic coherence of the two products is asserted more strongly than the Commercialization Plan supports. The Research Strategy argues Caregiver Staffing exists because capacity is required to establish care; the Commercialization Plan concedes that staffing revenue arises "regardless of where the underlying client originates... existing clients, externally generated referrals, growth, turnover replacement, or a capacity constraint observed through CareNavigator." On the applicant's own account, most staffing revenue will have nothing to do with CareNavigator, which makes the beachhead substantially a separate recruitment business rather than an integrated one. - Willingness-to-pay evidence for the beachhead is three providers. Three is enough to establish that the number is not zero; it is not enough to establish a price, a package, or a demand curve, and the Significance case leans on it. - Under the SBIR Significance criterion, the probability of commercialization is itself scored. That probability is materially reduced by a nil commercialization record across two prior Phase II awards, an IP position limited to trade secret and copyright, and a financing plan resting on non-commitments — and reviewers will import some of that into this score.

2. Investigator(s) — Score: 3

Assessment based solely on team information contained in the Research Strategy and Commercialization Plan; biosketches and letters assumed adequate.

Strengths - Genuine continuity rather than a team assembled for an application. Falohun was PD/PI on both prior NIA awards and, per the Progress Report, delivered against their milestones; DuBose was Co-I from Phase I through IIB; Fan holds the human-subjects protocols and served on the prior award; Ory has advised for more than six years and brings 20 years at NIA. This is the same group that built the thing being extended. - The disciplinary mix maps to the work: applied AI/software (Falohun), clinical and research operations (DuBose, a practicing primary-care clinician and CRO), epidemiology and mixed-methods evaluation with independent academic execution (Fan/Clemson), implementation and dissemination (Ory), and health-technology commercialization with an exit history and investor network (Qu, ~30 years). - Governance is explicit and workable: PI retains final go/no-go authority; Co-I owns the milestone calendar and research

operations; academic human-subjects work is executed by an independent university partner rather than by the company; independent statistical review and external CPA validation are named as checks. - The team has direct, sustained operating exposure to the failure modes it proposes to engineer around — two full-time family/provider call-center personnel for over three years, and 300+ provider discovery conversations.

Weaknesses - No member of the team has commercialized a product. The PI's commercialization record consists of ~\$5.7M in federal R&D and under 1% non-federal revenue. Qu supplies the missing experience, but as a quarterly advisor with an introduction role — not as an executive with operating authority over the commercial motion that Aim 3 is built to test. - The commercial function has no named owner. Aim 3 requires preregistering a priced offer, running structured pricing research, activating paid markets, and measuring conversion and retention; the Commercialization Plan assigns this to founders who already hold "overlapping responsibility for strategy, product, engineering, research, finance, and administration." - No named owner for agent safety, permissioning, or execution QA — the function most likely to determine whether the Aim 1 gate is passed honestly. For an application whose central innovation is bounded autonomous action on behalf of vulnerable elders, this is a conspicuous gap in the described structure. - Aim 3, Task 3.2 states that "execution by staff beyond the original founders will test whether market entry is transferable," but the Commercialization Plan defers hiring commercial and customer-success staff until "CRP milestones demonstrate repeatable provider sales" — i.e., until after the milestone that the non-founder staff are supposed to generate. The two documents describe a circular staffing sequence.

3. Innovation — Score: 3

Strengths - The framing avoids the most common failure mode in this space: the application does not claim novelty because it uses LLMs. It claims novelty in a *representation* — a longitudinal Care Establishment Model over seven eldercare domains carrying substates, timestamped events, geography, provenance, evidence, ownership of the next action, and explicit terminal states. That is a defensible and non-obvious engineering thesis, and the argument for why it is load-bearing ("software cannot reliably automate a fragmented pathway unless it can observe what state a case is in, what changed, who owns the next action, and what evidence defines completion") is the best paragraph in the application. - The measurement corollary is genuinely novel: directories describe what should exist and claims describe what was billed, but executed cases reveal what actually happened between recognized need and established care. Framing execution as an instrument rather than only as a service is a real contribution. - The field-learning concept is the most distinctive idea in the package and is well illustrated: a provider publicly accepts a payer but reports by phone that weekend capacity is unavailable until November; that fact improves the current case and every subsequent case in that market. The observation that eldercare operational truth is local, fast-changing, and largely absent from the public web is correct and is a real basis for a compounding data asset. - The bounded-execution loop (interpret, plan, permission, act, observe, update, wait or return, verify or escalate) with the explicit principle that "sending a form or leaving a voicemail is an event, not an outcome" is a mature and appropriately unglamorous statement of what the hard part actually is. - Human-control boundaries are drawn as design constraints rather than disclaimers: attestation, legal authority, submission, scheduling, and consequential family choice remain human; the system does not diagnose, prescribe, place orders, or make clinical decisions.

Weaknesses - The agent execution layer is the least defensible of the three innovations on a commercial time horizon. General-purpose agentic execution — browser automation, document handling, scheduling, voice — is being commoditized rapidly by platform vendors; the durable novelty is the domain state model and the local operational corpus, not the agents. The application does not distinguish which of its advantages survive a general-purpose agent becoming competent at portal navigation, and reviewers familiar with the trajectory will ask. - Innovation 3 is materially weaker than the other two. Recruiting pre-health students into paid caregiving is not novel — the applicant's own

Commercialization Plan names CareYaya as an existing student-caregiver model. The claimed innovation is the *data-directed targeting* of new supply at observed capacity deficits, but no part of the Approach tests whether deficit-targeted recruitment outperforms untargeted recruitment. Without that comparison, the innovation reduces to conventional student recruitment plus a dashboard. - The claim that the architecture is "labor-pool agnostic" is asserted and never tested. Only one pool (pre-health students) is used, and the Approach contains no task that would demonstrate transfer to a second pool. - Field learning has no accuracy criterion anywhere in the application (see Approach). An innovation whose stated value is knowledge quality is measured only on process compliance. - Nothing in the Innovation section is protected by anything stronger than trade secret, and the Commercialization Plan defers patent evaluation entirely. For innovations described as load-bearing and compounding, that is a thin position.

4. Approach — Score: 5

Strengths - The staged architecture is the application's best feature and is correctly ordered: technical verification precedes human exposure, real-world value precedes price, and payment does not relax any human-protection boundary established earlier. Aim 1 gates Aim 2; Aim 2 gates Aim 3. This is exactly right for a system that takes consequential actions. - The primary Aim 2 endpoint is well chosen and is the hardest available: the family must confirm that the primary aid or care need documented at intake was actually obtained within 90 days, with source verification where available, and the application states explicitly that "a referral, application, benefit approval, provider match, scheduled appointment, or job offer alone is not establishment." Most applicants in this space would have measured referrals. - Task 2.2's instrumentation is genuinely thorough, and the requirement that "a product failure is distinguishable from a funding restriction, provider refusal, workforce shortage, family choice, or loss to follow-up" is the correct discipline for interpreting a low establishment rate. - Aim 1's usability study is appropriately specified: 25 families, 25 providers, 25 student workers, moderated think-aloud on standardized non-live cases, SUS and 12-item TIAS with thresholds, semi-structured interviews with a shared codebook, independent coding, discrepancy resolution, audit trail, and a joint display translating findings into engineering changes. Using non-live cases before permitting real-world action is the right sequencing. - Task 3.1's preregistration of package, price range, assignment strategy, primary commercial outcome, follow-up horizon, and decision rules *before* the first paid market opens is unusually rigorous for a commercialization aim and directly prevents post-hoc redefinition of success. - Task 3.3's cost accounting is honest in ways applicants usually avoid: acquisition cost explicitly includes spend on accounts that do not convert, and cost to serve explicitly includes navigator time and exception handling — the two lines that most often hide a negative margin. - Willingness to pay is correctly demoted to "a pricing hypothesis, not commercial proof," and is elicited only after sufficient use. - Failure handling is credible and specific: prespecified alternatives proceed in sequence (correct value gaps, reduce cost, repackaging or reprice, retest), and the application commits to preserving and publishing evidence and stopping rather than scaling a failing model.

Weaknesses - **The primary Aim 2 success criterion does not exist in the application.** Table 2 states only that the day-90 proportion will be "estimated with prespecified precision," and the GO/NO-GO text conditions paid expansion on establishment "meet[ing] the prespecified success criterion" — a criterion that is never specified. Two of seven Aim 1 criteria are similarly empty ("prespecified completeness across representative pathways"). The application specifies thresholds precisely everywhere else (SUS ≥ 72 , TIAS ≥ 5 , $\geq 90\%$, $\geq 50\%$, $\geq 4.0/5$), which makes the omission at the single most consequential gate read as unresolved rather than merely unwritten. Reviewers cannot score a gate whose threshold will be chosen after review. - **N $\approx$ 400 is not justified.** It is described as a "target... subject to final power and recruitment justification before activation," with the final N "fixed by power/feasibility analysis." No power calculation, no target precision, no anticipated establishment rate, and no attrition allowance appear anywhere. At N=400 and $p\approx 0.5$ the 95% CI half-width is roughly ± 5 points; whether that is adequate cannot be assessed because the decision it informs is

unstated. - **The primary endpoint cannot detect whether the innovation did the work.** A family confirming establishment at day 90 is fully consistent with human navigators — Olera already employs two full-time call-center staff, and the Research Strategy repeatedly routes unsupported and ambiguous workflows to a human — executing every consequential step. Task 2.2 captures human interventions in the data, but no endpoint, criterion, or prespecified target in Table 2 measures the share of tasks executed autonomously, the rate at which families grant agent permissions, or the establishment rate by execution mode. Aim 2 could be fully successful and say nothing about Innovation 2. - **Aim 2 is single-arm and therefore cannot support the institutional value proposition.** No concurrent comparator, no usual-care arm, no waitlist, no stepped-wedge. Yet the Commercialization Plan's entire institutional engine rests on the claim that establishing care reduces downstream utilization. A staged multi-market rollout is already planned; a stepped-wedge or sequential-activation comparison was available at essentially no incremental cost and was not proposed. Reviewers will ask why. - **≥90% expected-action agreement is a weak bar for consequential autonomous action, and it is not adjudicated independently.** One action in ten diverging from expectation, before deployment against benefit applications for vulnerable elders, warrants stratification by severity — a trivial deviation and an incorrect submission are not the same event — and the criterion contains no such stratification. The Progress Report separately claims Aim 1 "verifies closed-loop execution against blinded expert review" and the Commercialization Plan says Aim 1 "independently verifies"; the Approach text describes only comparison against an internally constructed "technical test set" with no blinding or external adjudication anywhere. - **The field-learning criterion measures provenance, not correctness.** "Field observations with provenance and QA — 100% of persisted operational updates" can be fully satisfied by a system that learns and propagates false facts with impeccable metadata. No accuracy target, no independent verification sample, no staleness or decay policy, and no re-verification interval are specified for a knowledge layer whose whole premise is that local operational truth changes rapidly. - **The ≥50% screened-to-placement criterion is contradicted by the applicant's own data.** The Commercialization Plan reports ~900 applications, 100 accepted, 25 placed. If "screened" corresponds to "accepted," the observed rate is 25% and the criterion is double it, with no justification for the improvement. - **Family recruitment feasibility is asserted, not evidenced.** No funnel is given at any stage: visitors → account → completed assessment → eligible index episode → day-90 confirmation. National organic traffic is ~15,500/month spread across "nearly every county"; concentrating ~400 verified index episodes into a small number of counties within roughly a year is a different problem from national traffic, and paid acquisition is invoked without a single CAC figure. - **Aim 3's retention measurement is not achievable within the award.** Paid conversion begins in Year 3; retention is to be "reported at 3, 6, 9, and 12 months as available"; Task 3.3 asserts these measurements "will directly replace the... retention... assumptions used in the Commercialization Plan." Twelve-month retention will exist for approximately no customers. - **The market count is undetermined in the Approach itself.** Task 2.1 defers the number and boundaries of markets; Task 3.2 defers the number of paid markets to "the final experimental and revenue architecture in Aim 2." Since sample size, operating cost, provider density, worker supply, and revenue are all functions of market count, an unfixed market count leaves the reviewer unable to assess feasibility of anything downstream. - Aim 3, Task 3.4 does not produce institutional-buyer evidence; it produces the *specification* for a future study ("the population, endpoints, data linkage, comparison design, and follow-up required for a subsequent payer... proof-of-concept"), plus modeled avoided cost explicitly labeled non-causal. This is intellectually honest and commercially insufficient for what the Commercialization Plan monetizes in Year 4. - Table 3's "Independent financial validation — Delivered where commitment remains applicable" hedges a commitment the Commercialization Plan asserts flatly as "external CPA validation of the commercial unit-economics model."

5. Environment — Score: 4

Strengths

- The operating environment is real and in production, not proposed: a nationally deployed platform, an expert-curated database of 72,000+ provider and aid-program records across all fifty states with 578 program guides, 700+ providers who have claimed listings, and organic traffic grown 10× since 2023 at near-zero acquisition cost. - Four peer-reviewed evaluations with family caregivers, two of them cited with full journal references (JMIR Aging 2024; JMIR Form Res 2026), establish that this company completes and publishes its research obligations. - Human-subjects capacity is external and appropriate: Clemson University leads the academic effort, Fan holds every protocol, and Ory advises as she did on Phase IIB. The company is not self-reviewing its own research. - No regulatory dependency: the application states clearly that no FDA approval is required and no in-licensing is needed, removing two common sources of environment risk. - Compliance infrastructure is described concretely rather than gestured at: ADC accounting/CPA support, specialized legal and compliance counsel, HIPAA-aligned data practices, referral-compliance review, and federal research/IRB/data-stewardship experience. - The SBIR history section answers every required question directly and without evasion, including the unflattering answers.

Weaknesses - The commercialization track record is, by the applicant's own disclosure, effectively nil: SBIR/STTR-derived revenue of 100%, 100%, 100%, ~100%, and ~100% across FY2021–FY2025, with FY2025 one-off pilot fees under 1%, against ~\$5.7M in NIA support. The NIH Environment criterion asks directly how successful the company has been at commercializing prior SBIR work, and the honest answer here is "not yet at all." The applicant's explanation — that monetization was deliberately held until it could occur without charging families or gating referrals, and that provider willingness to pay requires geographic concentration research funding was not scoped to create — is coherent and partially persuasive, but it remains an explanation for the absence of a track record rather than a track record. - The team is very small for the proposed scope: founders, one full-time engineer of roughly two years, two full-time call-center personnel, part-time marketing, and two part-time research assistants. Year 1 alone requires building a seven-domain longitudinal state model, bounded agent execution with permission gates and scheduling, browser automation, document assembly, email/SMS/fax integration, a field-learning knowledge layer with QA, and staffing infrastructure through employer-confirmed placement — while running a 75-participant IRB study. Neither document contains an engineering hiring plan or effort allocation. - The company has apparently never raised private capital in six years of operation. Neither document discloses any equity financing, and 100% federal revenue implies none. This bears directly on the Environment question of whether the SBC can sustain and grow itself, and on the Fundraising Plan's credibility. - Prior NIA support is stated as running 2021–2027, overlapping the proposed three-year CRP. Neither document addresses concurrent-effort allocation across the overlapping period.

CRP Progress Report

Strengths

- The report is organized around risk retirement rather than activity, and the logic is genuinely cumulative: each stage tested one uncertainty before the dependent layer was built. Reviewers found this the most persuasive part of the package. - Every claim is tied to a specific funding source, which correctly separates what NIH paid for from what the company tested on its own capital — an unusually clean answer to the NOFO's request for progress "beyond the scope of the Phase II project." - The evidence is concrete and mostly verifiable: 72,000+ records and 578 guides; usability 4.57/5, acceptance 5.83/7 at four weeks (n=65), multi-agent 5.73/7 (n=31); ~50 → 500+ daily organic visits; 300+ provider discovery conversations and 700+ claimed listings; ~900 workforce applicants with 20+ placements; four providers trialed staffing and three paid. - The statement that manual delivery was deliberate — "to learn the workflow before building the infrastructure that automates it" — is exactly the right answer to the reviewer question of why a company

would hand-deliver a product it proposes to automate. - Third-party investment readiness is described without inflation, and the disclaimer "these relationships are not financing commitments" is stated twice. Reviewers credited the restraint. - The statement of what remains unretired is accurate and does not overclaim.

Weaknesses - The report contains the market-count statement — "Aim 2 measures established care and stakeholder value in **two real markets**" — that directly contradicts the Commercialization Plan's eight and Aim 2's own deferral. Of the three market counts in the package, this is the one embedded in the section reviewers read as the applicant's own summary of its plan. - It asserts that "Aim 1 verifies closed-loop execution against **blinded expert review**," a verification method that appears nowhere in Aim 1. This is the kind of unsupported strengthening that costs credibility precisely because the rest of the section is so disciplined. - The willingness-to-pay evidence is thinner than the surrounding prose implies: three paying providers, at \$50–\$275 per placement plus \$275/month, is a demonstration that the number is not zero. The Commercialization Plan then bases a \$250 base case on it. - "Placed more than 20" here versus "placed 25" in the Commercialization Plan is a small inconsistency in the package's single most load-bearing operating statistic. - The claim that the Phase IIB review "assessed its commercial potential as extremely high" quotes a prior summary statement. It is accurate but reviewers generally discount an applicant's citation of its own previous review, and it occupies space that a cost-per-placement or conversion figure would use better. - The cross-reference "the commercial endpoints in Tables 4 and 5 are the ones those investors said they would need to see" is broken: the Research Strategy has no Tables 4 or 5, and the Commercialization Plan's Tables 4 and 5 are the SBIR history and the market-acceptance hurdles, neither of which lists commercial endpoints.

Commercialization Plan

Strengths

- The Statement of Need is a correct and well-argued description of a real financing gap: investors must underwrite completion before institutional value is established, while institutional buyers require evidence before they will purchase. The applicant states specifically what would not otherwise happen, and why. - The rejection of the two obvious monetization shortcuts is principled and commercially reasoned, not merely ethical: charging families raises the barrier for households already unable to afford care, and charging providers for referrals introduces steering incentives and excludes some federally reimbursed providers. - The revenue model is arithmetically self-consistent throughout. $8 \text{ markets} \times 10 \text{ hires/month} \times \$250 = \$240\text{K}$ annualized; six paid-month equivalents = $\sim \$120\text{K}$ in Year 3; $15 \text{ markets} = \$450\text{K}$ in Year 4; $25 \text{ markets} = \$750\text{K}$ staffing plus $\$750\text{K}$ institutional = $\sim \$1.5\text{M}$ in Year 5. Every figure checks. - Table 7 is a model of disclosure discipline: each input is labeled as a published benchmark, an Olera assumption, or a CRP hypothesis, and the \$250K institutional figure is explicitly flagged "Planning hypothesis, not a benchmark." - The institutional engine is modeled conservatively and correctly: zero institutional revenue during the CRP, and evidence-gated contracts rather than an unvalidated PMPM multiplied across a hypothetical health plan. The application says so in as many words, and reviewers noticed. - Table 8's staged commercial framing (Build → Validate free → Monetize → Expand → Scale) matches the Research Strategy's gating logic, and Table 9's management timeline attaches a financing decision to each stage. - The county as the unit of commercialization is a genuinely useful construct: it makes acquisition cost, throughput, and care establishment measurable in the same denominator, and it supports a replication rather than market-share growth story across 3,100+ units. - Section 5's "temporal barriers to replication" argument is the right argument for this asset class — distribution, national data infrastructure, workforce relationships, and accumulated longitudinal records that a new entrant cannot obtain retrospectively.

Weaknesses - The single input that drives all staffing revenue — 10 successful hires per county per month, described as "conservative" — has no supporting evidence and is contradicted in magnitude by the applicant's own pilot. The pilot placed 25 workers total. The Year 3 base case requires 480 successful hires in six paid months across eight counties,

roughly 20× lifetime performance; the Year 3 exit run rate implies 960/year. At the pilot's ~2.8% applicant-to-placement yield, 480 hires requires on the order of 17,000 applicants in six months, or ~2,100 per county, against a pilot that produced ~900 in one market over multiple semesters. - The long-range illustration is arithmetically incompatible with the stated workforce wedge. Even at the *target* 50% screened-to-placement rate, 60,000 hires at 500 counties requires ~120,000 screened candidates — approximately 92% of the combined annual national MD (54,699) and NursingCAS (75,078) applicant pools the plan itself cites. At the observed pilot yield it requires ~2.1M applicants, roughly 16× those pools. The hedge that "these two pipelines are not the full addressable workforce" is then carrying the entire scale case. - There is no cost side. No cost to serve, no contribution margin, no gross margin, not even an illustrative range, anywhere in the plan. A "conservative base case" that models only revenue is not a base case. This matters acutely at \$250 per hire with a ~36-applicants-per-placement funnel: the plan's viability turns on a cost figure it does not contain. - No provider-conversion funnel is given. 700 providers have claimed listings and 3 have ever paid, but the plan never states how many paying accounts each market requires, what conversion rate is assumed from claimed to paid, or what repeat-purchase rate underlies "10 hires per month" from a given account. - The free-to-paid transition is unaddressed as a commercial risk. Staffing is deployed free to providers throughout Year 2 and priced in Year 3 to what appear to be the same markets; nothing addresses anchoring, or the possibility that Year 2 volume is a function of the zero price. - The institutional revenue line is not supported by the evidence the CRP produces. Year 4 books ~\$250K from ~1 relationship, priced against a single-arm establishment proportion, explicitly non-causal modeled avoided cost, and a study design for a proof-of-concept that has not been run. MA plan and ACO procurement cycles alone make a Year 4 contract of that size aggressive. - The IP position is weak against the criterion as written. Trade secret, copyright, and trademark, with patents only to be "evaluated," is a thin portfolio for a company whose defensibility argument is that its integrated asset is hard to recreate. No filings, applications, or provisional coverage are disclosed. - Section 7 states the CRP "applies this process across approximately eight markets," preempting a decision the Research Strategy explicitly defers, and conflicting with the Progress Report's two.

Fundraising Plan

Strengths

- The sequencing is sensible and non-naive: cultivation begins around the end of Year 2 on the strength of technical and early real-world evidence, with a formal raise supported in Year 3, explicitly to avoid meeting a funding gap at award end. - A specific, credentialed advisor is named with a defined cadence (quarterly), defined functions (investment-readiness milestones, pressure-testing the financing case, materials, introductions), and named target networks (Ziegler, Equitage Ventures, 7Wire Ventures, Alumni Ventures). - The raise is sized against a modeled gap (\$800K in Year 4 narrowing to ~\$350K in Year 5) with a stated rationale for why the ask (\$3–5M) exceeds the deficit — expansion capital, not runway. - A downside path is stated: moderate expansion and let revenue extend runway if financing is slow.

Weaknesses - Nothing in the plan constitutes a commitment, and the application says so — twice. The NIH criterion asks how well the application supports the ability to secure third-party funds equal to or exceeding the NIH request; \$3–5M against ~\$4M requested is at the boundary, supported only by advisory relationships. - No private financing history is disclosed. A company that has raised no equity in six years proposing a \$3–5M institutional round is a materially harder case than the plan acknowledges, and the plan does not address it. - The financing case rests on Year 3 metrics of ~\$120K recognized revenue and a ~\$240K run rate. That is a seed-stage traction profile being taken to a Series-A-sized raise, and the plan offers no argument for why those metrics clear the bar with the named funds. - Ziegler is an underwriter serving nonprofit senior living providers, not a venture investor in early-stage software. Naming it alongside 7Wire and Alumni Ventures without distinguishing the roles suggests relationship breadth substituting for financing-path specificity.

Project Management Plan

Strengths - Decision authority is unambiguous: the PI holds final go/no-go at major decision points. - Named workstream owners, maintained task boards, regular operating meetings, and milestone dashboards tied explicitly to both the Research Strategy and the Commercialization Plan. - The governing rule is stated correctly: "a milestone that misses its predefined threshold triggers the corresponding alternative strategy rather than automatic continuation." - External checks are built in — independent statistical review and external CPA validation of the unit-economics model. - Table 9 attaches a commercialization decision *and* a financing gate to each project year, which is what a CRP management plan should do.

Weaknesses - The plan's own milestones are disclaimed. Table 9's footnote states that "revenue values are planning projections rather than CRP success criteria," which removes accountability from the only quantified milestones in the table and leaves Years 3–5 with decision gates but no pass/fail thresholds. - No workstream owners are actually named — the mechanism is described, the people are not, and the only named individuals are the two founders, an academic Co-I, and advisors. - No effort allocation, FTE plan, or hiring schedule appears in either document, for a program whose Year 1 engineering load substantially exceeds the described engineering staff. - The plan contains no schedule-contingency logic. The Research Strategy notes contingency time is retained; the management plan does not say what is descoped, or by whom, if Aim 1 verification slips and compresses Aims 2 and 3. - External CPA validation is asserted here but hedged in Research Strategy Table 3 ("delivered where commitment remains applicable").

Study Timeline

Strengths

- The work is genuinely staged rather than merely sequenced, with downstream human and commercial exposure explicitly conditioned on upstream gates, and parallelization only where it is safe. - Contingency time for remediation and re-verification is explicitly retained. - A task-level three-year timetable is provided (Figure 9), and the commercialization management timeline in Table 9 is period-by-period.

Weaknesses - Year 1 is heavily overloaded: four build-and-verify tasks spanning state modeling, agentic execution across six-plus tool modalities, a QA'd field-learning layer, and staffing infrastructure, *plus* a 75-participant IRB study with think-aloud sessions, SUS/TIAS instruments, and qualitative coding — all before any real-world execution is permitted. - Aim 2 must complete enrollment of ~400 index episodes, allow 90 days per episode, run post-use surveys and a purposive interview sample spanning four pathway outcomes, complete a joint-display integration, and pass a GO/NO-GO before Aim 3 opens. Enrollment must therefore close four to five months before Aim 2 ends, which is not acknowledged. - Aim 3 is compressed into what remains — preregistration, structured pricing research, activation of paid markets, six paid-month equivalents of revenue, live-record economics with external validation, actuarial modeling, and the institutional evidence package. - Retention at 9 and 12 months is not observable within the award for customers converting in Year 3. - The timeline cannot be fully assessed because market count is unfixed; activating two counties and activating eight are materially different schedules.

Human Subjects Design Visible in the Research Strategy

Note: Full Human Subjects and Clinical Trials materials were not supplied and are assumed adequate for this mock review, per instruction. No determination is issued on their basis. The findings below concern only design features intrinsic to the Research Strategy that no separate Human Subjects section could cure. **Assessment: CONCERN.**

Strengths - Human protection is architectural rather than procedural. Aim 2 opens only after prespecified technical criteria are met, required approvals cannot be bypassed, users demonstrate acceptable usability and control, and no critical human-protection issue remains unresolved. - "Human protection is overriding" is stated as a standing rule, and any unresolved critical privacy, permission, unauthorized-action, or safety failure is a NO-GO for real-world execution. - The scope limitation is genuinely protective and clearly drawn: the system does not diagnose, prescribe, place clinical orders, or make clinical decisions, and authorized health information may be used only to extract and execute administrative tasks already defined by the family, provider, or clinician. - Consequential actions requiring attestation, legal authority, submission, scheduling, or family choice require human approval; unsupported or ambiguous workflows remain human-assisted; the family remains the decision-maker. - Non-live standardized cases are used for the Aim 1 usability study, so no participant is exposed to agent action before verification. - A pause-and-review mechanism is prespecified for critical privacy events, unauthorized consequential actions, systematic erroneous execution, and serious safety signals, with explicit provision for required IRB action and re-verification. - Human-protection boundaries are stated to be unchanged by the introduction of payment in Aim 3.

Weaknesses - The $\geq 90\%$ action-agreement gate is not stratified by severity. A 10% rate of trivial deviations and a 10% rate of erroneous submissions to a state LTSS program are entirely different risk profiles, and only the latter can harm a participant. Severity stratification belongs in the Approach, not in a Human Subjects form. - "Predefined serious safety signal" is invoked without any adverse-event taxonomy. For an intervention whose harms are administrative rather than clinical — a denied or delayed benefit, a lost eligibility window, an erroneous disclosure to a provider or agency — the harm categories are non-obvious and need to be defined where the study is designed. - No independent safety monitoring body is named. This study lets software take consequential real-world actions affecting benefits access for a vulnerable population; independent monitoring is warranted on design grounds regardless of whether the work meets the clinical-trial definition. - Decisional capacity and surrogate authorization are not addressed anywhere in the Research Strategy, despite the fact that the mechanism is authorization: a family "authorizes" the system to prepare an LTSS application or contact agencies. Who may grant that authorization, how capacity is assessed, and what happens when the care recipient and the authorizing family member disagree are design decisions intrinsic to the intervention. *The procedural elements of consent are assumed adequately addressed in the full Human Subjects materials; the authorization model itself is not.* - "Permissioned browser automation" against government and provider portals is proposed without addressing portal terms of use or the risk that automated access jeopardizes a participant's own account or application status. This is a participant-risk question, not only a legal one. - Aim 2 sets no ceiling on human-protection incidents below the "critical" threshold. A workflow producing a steady rate of non-critical permission or execution errors has no defined stopping rule.

SUPPLEMENTARY ANALYSES

A. Cross-Document Consistency Findings

#	Element	Research Strategy	Commercialization Plan	Classification
1	Number of markets	Progress Report: "two real markets." Task 2.1: number "will be finalized" later. Task 3.2: paid market count "fixed from the final... architecture in Aim 2."	§7: "approximately eight markets." Table 8: ~8 free (Y2), ~8 paid (Y3), ~15 (Y4), ~25 (Y5).	Potential score-driving contradiction
2	Aim 1 verification method	Task 1.2: comparison against internally constructed technical test set; no blinding, no external adjudicator. Progress Report: "verifies closed-loop execution against blinded expert review."	§1: Aim 1 "independently verifies the execution and outcomes technology."	Potential score-driving contradiction
3	Y3 paid markets: same or new?	Task 3.2: "Enter new paid markets and test conversion and replication."	Table 8 shows ~8 paid in Y3, not 8+8, implying the Y2 free markets convert.	Material inconsistency — under the RS reading the Y2 free markets never generate revenue; under the CP reading replication is never tested in a fresh market.
4	Institutional evidence produced	Task 3.4 delivers a <i>specification</i> for a future payer proof-of-concept plus explicitly non-causal modeled avoided cost.	Y4 books ~\$250K institutional revenue from ~1 relationship.	Material inconsistency — CP monetizes evidence RS does not produce.
5	Workforce placements in pilot	"placed more than 20 into provider jobs."	"accepted 100 candidates, and placed 25 with local providers."	Minor inconsistency (but in the package's most load-bearing operating statistic).
6	Observed price vs. modeled price	Progress Report: three providers paid "\$50 to \$275 per placement and \$275 a month."	Base case \$250 per successful hire; sensitivity \$150–\$350; no monthly component modeled.	Material inconsistency — base case sits at the top of a three-observation range and drops the recurring form actually observed.
7	Screened → placement rate	Aim 1 criterion: ≥50%.	Pilot: 100 accepted → 25 placed = 25%.	Material inconsistency — criterion is 2× the only observation, unjustified.
8	Independent financial validation	Table 3: "Delivered where commitment remains applicable."	§9: "external CPA validation of the commercial unit-economics model" asserted as a standing check.	Minor inconsistency.
9	Posture of Aim 3	Aim 3 tests whether value converts; scale is a GO/NO-GO <i>outcome</i> .	§1: "Aim 3 scales deployment and builds the institutional-buyer evidence case."	Minor inconsistency.
10	CRP component decomposition	Three innovations: Care Establishment Model; bounded agent execution + field learning; data-directed workforce creation.	Three components: Caregiver Staffing; Task-based AI Agent Execution; Analytic Outcomes Data layer.	Minor inconsistency — two different tripartite decompositions of the same work; the Care Establishment Model, the RS's load-bearing innovation, has no named counterpart in the CP.

#	Element	Research Strategy	Commercialization Plan	Classification
11	Table/figure cross-references	Two Tables numbered 1 and two numbered 2, then a jump to Table 6; Progress Report cites "Tables 4 and 5," which do not exist in the RS and do not contain endpoints in the CP.	Tables 1–9 numbered cleanly.	Minor inconsistency (presentation risk).
12	Traffic figures	"more than 500 [visits] a day."	"15,500+ visitors per month"; "more than 500 per day in 2026."	Aligned.
13	Customer/user/buyer definitions	Table 1: families free users; providers free participants and paid staffing customers; risk-bearers evidence-gated.	S4: identical structure.	Aligned.
14	Employment liability boundary	Olera screens and refers; licensed providers retain hiring, training, employment, supervision.	Identical.	Aligned.
15	Free-to-family principle	Payment does not gate the care-establishment pathway.	Core access free; no referral gating.	Aligned.
16	Revenue arithmetic	n/a	$8 \times 10 \times \$250 = \$240K$; six months = \$120K; 15 markets = \$450K; 25 = \$750K + \$750K = \$1.5M.	Aligned — every figure verifies.

B. Claim–Evidence Skepticism Test

#	Claim	Evidence in supplied documents	Strength	Remaining uncertainty	Score relevance
1	Providers will pay ~\$250 per successful hire	Three providers paid \$50–\$275/placement and \$275/month in a manual pilot	Weak	Whether \$250 survives at volume, under a non-founder sales motion, in unfamiliar counties; the base case is at the 90th percentile of an n=3 sample and drops the recurring component actually observed	Significance, Comm. Plan, Overall
2	10 successful hires per county per month is conservative	None. Pilot placed 25 total	Unsupported	The entire revenue model. Y3 requires 480 hires in six months (~20× lifetime)	Comm. Plan, Approach, Overall
3	≥50% of screened workers reach employer-confirmed placement	Pilot: 100 accepted → 25 placed (25%)	Weak / contradicted	Why the rate doubles; whether "screened" and "accepted" denote the same stage	Approach
4	CareNavigator increases verified care establishment within 90 days	None yet; Aim 2 designed to measure it	Unsupported by design, appropriately	No comparator, no prespecified threshold, no automation-share endpoint — so even a successful Aim 2 leaves attribution open	Approach, Significance
5	Bounded agent execution can be verified before real-world use	Phase IIB multi-agent LLM in production; acceptance 5.73/7 (n=31)	Moderate for the reasoning layer; Weak for execution	No preliminary data of any kind on browser automation, fax, portal interaction, document assembly, or voice	Innovation, Approach

#	Claim	Evidence in supplied documents	Strength	Remaining uncertainty	Score relevance
6	Field learning improves routing and future cases	Illustrative example only; criterion measures provenance, not accuracy	Unsupported	Whether learned facts are correct, how staleness is handled, whether reuse improves outcomes	Innovation, Approach
7	Olera can reach families at meaningful scale	50 → 500+ daily organic visits, 15,500+/month, "nearly every county," near-zero CAC	Moderate	National traffic ≠ county-level density; no funnel from visit to index episode; no family CAC figure	Significance, Approach
8	Institutional buyers will contract at ~\$250K beginning post-CRP Y4	GUIDE Model precedent; 35.2M MA enrollees; 511 MSSP ACOs / 12.6M beneficiaries; one AHRQ association study	Weak	Precedent establishes the buyer class, not Olera's product, price, or procurement path; Task 3.4 yields a study design, not payer-grade evidence	Significance, Comm. Plan, Overall
9	Olera can raise \$3–5M post-CRP	Named advisor and networks, explicitly "not financing commitments"; no prior equity raise disclosed; ~100% federal revenue for five years	Weak	Whether ~\$120K revenue and a ~\$240K run rate clear an institutional round; no term sheets, no prior raise, no valuation logic	Fundraising Plan, Environment, Overall
10	The integrated asset is defensible	Trade secret, copyright, trademark; patents "will evaluate"; accumulating-data argument	Weak on formal IP; Moderate on the temporal/data argument	What is actually protectable if a general-purpose agent platform commoditizes the execution layer	Comm. Plan

Two further claims warrant flagging without full treatment: "Aim 1 verifies closed-loop execution against blinded expert review" (Progress Report) is *Unsupported by the Approach* text; and "the CRP applies this process across approximately eight markets" (CP §7) is *Unsupported by a Research Strategy* that twice defers the number.

C. Commercial-Assumption Retirement Test

Assumption	Evidence today	Aim / Task	Endpoint	Success criterion	Truly retired?	Residual uncertainty after CRP
Providers will pay a defined price for placements	3 paying providers, manual delivery	3.1, 3.2	Paid conversion by offer/market	"Prespecified before rollout" — not stated	Partially. Tests price acceptance in ~8 markets over ~6 months	Whether price holds at scale, in unfamiliar geographies, without founder involvement
Throughput of ~10 hires/county/month	None	2.1–2.3 (free deployment)	Hires per county per month	Not stated as a criterion	Partially — and this is the right design. Y2 measures throughput <i>before</i> price is introduced	Throughput under a <i>paid</i> regime; provider demand elasticity is untested until Y3
Applicant supply is reachable and replenishable in new counties	~900 applicants, one market	2.1, 3.2	Applicant acquisition and CAC by market	None specified	Partially	Whether university relationships exist in target counties; whether yield holds outside the home market

Assumption	Evidence today	Aim / Task	Endpoint	Success criterion	Truly retired?	Residual uncertainty after CRP
Screened workers convert to placements at $\geq 50\%$	25% observed	1.4	Employer-confirmed placement rate	$\geq 50\%$	Yes, mechanically — but the criterion is unjustified and may be set above achievable	If the gate fails, no alternative is specified
Verified care establishment occurs at a commercially meaningful rate	None	2.3	Day-90 verified establishment proportion	Not stated	No. A gate with an unstated threshold cannot retire an assumption	Whether the observed rate is good, and against what
Established care creates institutional economic value	Association studies only (52% vs 34% ED use)	3.4	Modeled avoided utilization, explicitly non-causal	Package "delivered with causal boundaries explicit"	No — by the applicant's own statement	Essentially total. Yet Y4/Y5 institutional revenue depends on it
Unit economics support a business	None	3.3	CAC, cost to serve, margin, retention from live records	"Measured"	Substantially yes — the strongest retirement in the package	One price band, ~8 markets, ~6 months; 9- and 12-month retention unobservable
Market entry is transferable beyond the founders	None	3.2	Playbook performance and deviations	"Reported"	Weakly. Reported, not thresholded	Compounded by the CP's deferral of commercial hiring until after the milestone this tests
Families will authorize AI agents to act	Acceptance 5.73/7 (n=31) on advisory use	1.5, then 2.2	SUS ≥ 72 , TIAS ≥ 5 on non-live cases	Stated for Aim 1	Weakly. Trust on hypothetical cases $\neq$ authorization behavior on a real LTSS application; no Aim 2 endpoint measures permission-grant rate	The single most important adoption metric for Innovation 2 is not measured
The asset is defensible	Trade secret posture	None	None	None	No — untested by any aim	Whether anything survives commoditization of the execution layer

D. Feasibility Stress Test

Dimension	Verdict	Basis
Family acquisition (~400 index episodes)	Plausible but vulnerable	Real national demand (15,500+ monthly visitors, near-zero CAC) but no funnel at any stage, no family CAC, no attrition allowance, and national traffic must be concentrated into a small number of counties with 90-day confirmed follow-up
Provider acquisition (participation)	Credible	700+ claimed listings, an existing profile-claim flow, 300+ discovery conversations, and franchise relationships
Provider acquisition (paid conversion volume)	Weakly supported	Three paying customers in history; no stated conversion rate from claimed to paid; no stated number of paying accounts needed per market
Workforce acquisition — CRP period	Plausible but vulnerable	480 hires in six months against 25 lifetime placements; at pilot yield this needs ~17,000 applicants, ~2,100 per county

Dimension	Verdict	Basis
Workforce acquisition — post-CRP scale illustration	Implausible as stated	500 counties / 60,000 hires needs ~120,000 screened at the <i>target</i> rate — ~92% of the combined annual national MD + NursingCAS applicant pools the plan itself cites; ~2.1M applicants at observed yield
Market activation	Weakly supported	Not because eight counties is inherently unreasonable, but because the number is contradicted across documents (2 / ~8 / TBD) and no activation cost or staffing model is given for running eight simultaneously
Technical workload	Plausible but vulnerable	Seven-domain state model + agents + browser automation + document assembly + email/SMS/fax + field-learning QA + staffing infrastructure in Year 1, with one full-time engineer plus founders and no hiring plan in either document
Human-subjects workload	Credible	Clemson executes; Fan holds protocols; Ory advises; the 75-participant Aim 1 study is routine for this team, which has completed four published evaluations
Commercial workload	Plausible but vulnerable	Preregistration, pricing research, paid market activation, six paid-month equivalents, live-record economics with external CPA validation, actuarial modeling, and the institutional evidence package all inside Year 3

E. Top Score-Driving Strengths

1. A first-rate Significance section: urgent, well-cited unmet need with a clean and consistently maintained separation of user, participant, and buyer, and explicit acknowledgment that the larger market is not yet purchasable.
2. A real technical thesis — making the care-establishment pathway computable — that is argued rather than asserted, and that correctly identifies representation as the prerequisite for both automation and measurement.
3. A gated architecture in which technical verification precedes human exposure, real-world value precedes price, and human protection overrides everything, with an explicit commitment to publish a negative result and stop.
4. A cumulative, source-attributed, honestly-bounded risk-retirement history: national deployment, four peer-reviewed evaluations, 700+ claimed provider listings, ~900 workforce applicants, 25 placements, and three providers who paid real money.
5. An endpoint that is genuinely hard and genuinely right — family-confirmed establishment of the documented need within 90 days, with referrals, approvals, matches, and job offers explicitly excluded.
6. Unusual disclosure discipline in the revenue model: hypotheses labeled as hypotheses, benchmarks labeled as benchmarks, zero institutional revenue assumed during the award, arithmetic that verifies throughout, and CAC defined to include non-converting spend.

F. Top Score-Driving Weaknesses

1. The primary Aim 2 success criterion and the N justification are absent; a gate whose threshold is chosen after review cannot be scored.
2. No endpoint distinguishes agent-executed from human-executed pathways, so Aim 2 can succeed without validating the central innovation.
3. The staffing throughput assumption driving all revenue is unsupported and ~20× demonstrated performance, and the long-range illustration exceeds the applicant's own cited labor pools.
4. Market count differs materially between the two documents (two vs. approximately eight vs. undetermined), destabilizing sample size, cost, and revenue simultaneously.
5. Aim 2 is single-arm, so the institutional value proposition the Commercialization Plan monetizes in Years 4–5 cannot be causally supported by the evidence the Research Strategy produces — a gap the Research Strategy concedes and the Commercialization Plan does not.

6. No cost side anywhere in the Commercialization Plan: no cost to serve, no margin, no illustrative range.
7. Weak formal IP and a Fundraising Plan resting on relationships the applicant twice describes as non-commitments, from a company with no disclosed private financing history and ~100% federal revenue for five consecutive years.
8. A very lean team against a Year 1 engineering load, with no hiring plan, no effort allocation, and a circular sequence in which non-founder staff are needed to produce the milestone that is supposed to justify hiring them.

G. Likely Discussion Dynamics

- **The market-count discrepancy will open discussion** and will be raised as a proxy for whether the applicant has actually converged on a plan. Because sample size, cost, revenue, and workforce demand all depend on it, reviewers will treat two-versus-eight as evidence that the operating plan and the financial model were written against different assumptions. The applicant's defense — that market count is properly derived from a power analysis — is sound but is undermined by the Commercialization Plan having already fixed it at eight and built a revenue model on it.
- **The strongest advocate's argument** will be that this is the correct use of the CRP mechanism: a real, deployed, published, used product with no commercial proof, a company that has been honest about exactly which risks remain, and a gated plan that will produce a clean answer either way. That advocate will emphasize that the applicant proposes to measure whether care is *established*, not referred — a standard almost no competitor holds itself to.
- **The skeptic's argument** will be arithmetic, and it will be delivered in two moves: 25 lifetime placements versus 480 required in six paid months; and 60,000 hires requiring ~92% of the national MD and nursing applicant pools at the applicant's own target conversion. The skeptic will add that ~\$5.7M of federal support has produced under 1% non-federal revenue and that this application asks for ~\$4M more.
- **The pivotal exchange** will be over whether the missing Aim 2 threshold is a drafting defect or a substantive one. Advocates will call it correctable; skeptics will note that the same document specifies SUS ≥ 72 , TIAS ≥ 5 , $\geq 90\%$, $\geq 50\%$, and $\geq 4.0/5$ with precision, which makes the omission at the decisive gate look like unresolved thinking. This argument determines whether Approach is scored 4 or 5 and therefore whether the overall lands at 3 or 4.
- **A methodologically strong reviewer will raise the single-arm design** and will point out that a stepped-wedge or staggered-activation comparison was available essentially for free given that markets are already being activated sequentially. This is the criticism most likely to change minds, because it identifies a fixable design choice that would materially increase the value of everything downstream.
- **Expect a Significance split of 2 versus 3**, turning on how much of the commercialization-probability question reviewers import into Significance versus quarantine in Environment and the Commercialization Plan assessment.
- **Expect the automation-share point to be raised late** and to land hard, because once articulated it is unanswerable from the document: nothing in Aim 2 prevents a successful result produced entirely by human navigators.

H. Reader-Experience Findings

Reported only for issues a competent Specific Aims page could not resolve.

- **Primary reviewer's mental model.** The Research Strategy builds a clear and compelling one: a fragmented pathway, an instrument that makes it computable, agents that execute within it, workforce that supplies its missing input, and three aims that build, prove, and price. The model holds through Innovation and begins to loosen in Approach at the first empty threshold, and breaks on contact with the Commercialization Plan's market count.
- **Most memorable issues for the secondary reviewer:** (1) "twenty-five placements, four hundred and eighty required"; (2) "two markets or eight?"; (3) "what is the establishment rate that counts as success?"; (4) "the plan needs causal evidence the aims don't produce"; (5) "one full-time engineer."

- **Where the skeptical reviewer's belief weakens:** at Commercialization Plan Table 7, on the line "Successful hires per county per month — 10 — Conservative Olera and CRP hypothesis." A number labeled conservative that is twenty times demonstrated performance re-frames every prior claim as potentially similarly calibrated, and the reviewer re-reads backward from there. A secondary break occurs at the Progress Report's "blinded expert review," which does not appear in the aim it describes.
 - **Structural presentation problems that survive a strong Specific Aims page:** the Research Strategy contains two Tables numbered 1 and two numbered 2 before jumping to Table 6, and the Progress Report cites "Tables 4 and 5" that exist in neither document with the described content. Individually trivial; in aggregate they cue a reviewer to check other cross-references, which is exactly when the market-count discrepancy is found.
 - **A structural asymmetry the Aims page cannot fix:** the Research Strategy's tripartite innovation (Care Establishment Model / agentic execution and field learning / data-directed workforce creation) does not map onto the Commercialization Plan's tripartite product (Caregiver Staffing / Task-based AI Agent Execution / Analytic Outcomes Data). The Care Establishment Model — the Research Strategy's declared load-bearing advance — has no named counterpart in the Commercialization Plan at all, so the reviewer who reads the commercial document first never learns what the technical core is.
-

FINAL SCORECARD

Criterion	Score	Primary reason
Significance	2	Urgent, well-documented need; unusually clean user/participant/buyer architecture; honest competitive analysis; institutional market correctly labeled evidence-gated. Held off 1 by an unsized beachhead market, willingness-to-pay evidence of n=3, a synergy claim the plan itself partly retracts, and commercialization probability drag
Investigator(s)	3	Genuine continuity, delivery against two prior NIA awards, complementary technical/clinical/research/commercialization expertise, clear go/no-go authority. Held off 2 by no commercialization success by any team member, no named commercial or agent-safety owner, and a circular hiring sequence
Innovation	3	The Care Establishment Model and field-learning concepts are substantive and well argued; execution-as-instrument is a real contribution. Held off 2 because the agent layer is commoditizing, Innovation 3 is conventional recruitment plus targeting that is never tested against untargeted recruitment, and "labor-pool agnostic" is untested
Approach	5	Excellent gating architecture, a hard and correct primary endpoint, rigorous cost accounting, and preregistration of the paid offer — undercut by an unstated primary success criterion, an unjustified N, an endpoint blind to automation share, a single-arm design that cannot support the institutional claim, a provenance-not-accuracy field-learning criterion, a placement criterion contradicted by the applicant's own data, and unsupported recruitment arithmetic
Environment	4	Real national deployment, published evaluations, external IRB capacity, no regulatory dependency, concrete compliance infrastructure. Held down by a nil SBIR commercialization record across ~\$5.7M, a very lean team against the Year 1 load, and no private financing history

Additional Assessments

Item	Assessment
Human Subjects design visible in Research Strategy	CONCERN (<i>full Human Subjects materials assumed adequate and not reviewed; no determination issued</i>)
Study Timeline	CONCERN
Commercialization Plan	MODERATE
Fundraising Plan	WEAK
Project Management Plan	MODERATE

OVERALL IMPACT SCORE: 4 (Very Good)

Why this score. The application targets a large, urgent, well-documented problem with a technically substantive and non-obvious thesis, proposes a correctly gated program that would produce an interpretable answer in either direction, and comes from a team that has delivered against two prior NIA awards and has real if small commercial traction. It is also the right project for this mechanism: the applicant is demonstrably past feasibility and demonstrably short of commercial proof, and says so plainly. Against that, the application's decisive quantitative commitments are either missing (the Aim 2 success criterion, the power justification, cost to serve), unsupported (staffing throughput, the $\geq 50\%$

placement rate, family recruitment yield), or inconsistent between the two documents (market count, verification method, the institutional evidence the plan monetizes). A 4 reflects a proposal whose architecture would likely support a sustained influence on the market offering, but whose execution parameters a panel cannot currently verify.

Why not one point better (3). A 3 would require the reviewer to believe the CRP will produce a decisive commercial answer. It cannot yet, for three specific reasons. The primary gate has no threshold, so "success" in Aim 2 is undefined at the time of review. The primary endpoint does not distinguish the proposed innovation from the human process it replaces, so a positive result would be ambiguous about the product. And the Commercialization Plan's larger revenue engine depends on causal evidence the Research Strategy explicitly declines to produce, so the most valuable commercial question survives the award unretired. Any one of these might be absorbed; together they mean the application's own logic does not guarantee the answer it promises.

Why not one point worse (5). A 5 would imply the strengths are ordinary. They are not. The customer architecture is clearer than most applications in this space achieve; the endpoint standard — family-confirmed establishment, with referrals and approvals explicitly excluded — is harder than the field's convention and is the correct standard; the risk history is cumulative, source-attributed, and honest about what remains; the cost accounting deliberately includes the two lines applicants usually hide; the pricing offer is preregistered before conversion is observed; and the applicant commits in writing to publishing a negative result and stopping rather than scaling a failing model. The revenue model, whatever the merits of its central input, is arithmetically consistent and explicitly labels its hypotheses as hypotheses. The identified weaknesses are specific and largely correctable within the existing design rather than reflecting a misconceived project.

FINAL APPLICANT DIAGNOSTIC

Prepared after scores were locked. Recommendations are limited to changes within the Research Strategy and Commercialization Plan; no recommendation depends on adding materials outside them.

Must fix in the Research Strategy before submission

1. State the Aim 2 primary success criterion and its number. - *Section:* Aim 2, Task 2.3 and Table 2; GO/NO-GO 2. - *Problem:* The primary criterion is "prespecified precision" against an unstated success threshold. - *Reviewer inference:* The applicant has not decided what counts as success, or is preserving the option to decide after seeing the data — which is precisely what the Aim 3 preregistration language elsewhere disavows. - *Scoring consequence:* This alone is worth a full point in Approach and is the swing factor between an overall 3 and 4. - *Minimum change:* One sentence with three numbers: the null/benchmark establishment rate the applicant regards as the status quo, the rate the CRP must exceed, and the CI half-width at the proposed N. E.g. "GO requires the lower bound of the 95% CI for day-90 verified establishment to exceed X%."

2. Justify N≈400 with an actual power/precision calculation and an attrition allowance. - *Section:* Task 2.3; Table 2. - *Problem:* N is a target "subject to final power and recruitment justification," with no calculation and no allowance for loss to follow-up on a 90-day family-confirmed endpoint. - *Reviewer inference:* The number was chosen for operational convenience and will be revised downward under recruitment pressure. - *Scoring consequence:* Approach; compounds the feasibility concern. - *Minimum change:* Three sentences — assumed establishment rate, target CI half-width, resulting analyzable N, plus an enrollment target inflated for stated expected attrition.

3. Add an endpoint that measures automation share. - *Section:* Task 2.2 / Task 2.3 and Table 2. - *Problem:* Nothing distinguishes agent-executed from human-executed pathways, and no measure of permission-grant rate exists. - *Reviewer inference:* Aim 2 may succeed on the strength of two call-center employees, validating nothing about Innovation 2. - *Scoring consequence:* Approach and Innovation; this is the criticism most likely to be raised last and answered worst. - *Minimum change:* Add two rows to Table 2 — (a) proportion of eligible consequential tasks executed autonomously within the verified envelope, with a target; (b) proportion of enrolled families granting agent authorization, with a target — and report establishment rate stratified by execution mode.

4. Add a concurrent comparison, or state explicitly why none is possible. - *Section:* Aim 2 analysis; Task 2.1 market activation. - *Problem:* A single-arm day-90 proportion cannot support the institutional value proposition the Commercialization Plan sells. - *Reviewer inference:* The applicant has not confronted what a payer will actually require. - *Scoring consequence:* Approach, and indirectly Significance and the Commercialization Plan assessment. - *Minimum change:* Since markets are already activated in sequence, convert the rollout to a staggered/stepped-wedge activation and analyze establishment against pre-activation periods in the same counties. If genuinely infeasible, say so in one sentence and state what comparison Task 3.4's design specifies instead.

5. Stratify the ≥90% agent-action criterion by severity and describe independent adjudication. - *Section:* Task 1.2; Aim 1 Table 1. - *Problem:* A single 90% threshold treats a formatting deviation and an erroneous benefits submission identically, and the Progress Report claims blinded expert review that Task 1.2 does not describe. - *Reviewer inference:* The gate protecting vulnerable participants from consequential machine action is weaker than the document elsewhere implies. - *Scoring consequence:* Approach; drives the Human Subjects design assessment from Concern toward Adequate if fixed. - *Minimum change:* Split the criterion — e.g. ≥90% overall action agreement *and* zero critical-severity divergences *and* zero bypassed approvals — define the severity classes in one sentence, and either describe the blinded independent adjudication or delete the claim from the Progress Report.

6. Give the field-learning layer an accuracy criterion. - *Section:* Task 1.3; Aim 1 Table 1. - *Problem:* "100% of persisted operational updates" measures metadata compliance, not truth. - *Reviewer inference:* The applicant is measuring what is easy rather than what matters, in the innovation it describes as most distinctive. - *Scoring consequence:* Approach and Innovation. - *Minimum change:* Add one row — independently verified accuracy on a random sample of persisted observations, with a threshold — plus one sentence on staleness/re-verification interval.

7. Justify the $\geq 50\%$ screened-to-placement criterion against the pilot's 25%. - *Section:* Task 1.4; Aim 1 Table 1. - *Problem:* The criterion is double the only observed rate, unexplained. - *Reviewer inference:* Criteria are being set to what is needed rather than to what is achievable. - *Scoring consequence:* Approach; feeds the credibility question about throughput. - *Minimum change:* Define "screened" precisely, state the pilot's rate at that same stage, and give the specific infrastructure change (verification before handoff, availability matching) expected to produce the improvement.

8. Resolve the market count and make it consistent with the Commercialization Plan. - *Section:* Progress Report ("two real markets"); Task 2.1; Task 3.2. - *Problem:* Three different answers across the package. - *Reviewer inference:* The research plan and the financial model were built on different assumptions. - *Scoring consequence:* Approach and Commercialization Plan; this is the finding most likely to open discussion. - *Minimum change:* Commit to a number or a tight range, state it identically in both documents, and if it must remain conditional, state the decision rule and the bounds ("no fewer than X, no more than Y, fixed by [date] using [rule]").

Must fix in the Commercialization Plan before submission

9. Support or restate the 10-hires-per-county-per-month assumption. - *Section:* §8, Table 7 and the surrounding projection. - *Problem:* The input driving all revenue is unsupported and is labeled conservative while exceeding lifetime performance by roughly 20×. - *Reviewer inference:* If this number is calibrated this way, every other number may be. - *Scoring consequence:* Commercialization Plan rating, Significance, and Overall Impact. This is the single highest-yield fix in the package. - *Minimum change:* Show the funnel — applicants per county per month, screening yield, provider interview rate, hire rate — with the pilot's actual values in a parallel column, and either derive 10 from that funnel or lower the base case to what the funnel supports and present 10 as the mature-state case. A defensible 4 is worth more than an undefended 10.

10. Add a cost side. - *Section:* §8. - *Problem:* No cost to serve, no contribution margin, no gross margin anywhere. - *Reviewer inference:* At \$250 per hire with a ~36-applicant funnel, the applicant may not know whether the unit is profitable. - *Scoring consequence:* Commercialization Plan rating and Fundraising Plan credibility. - *Minimum change:* One row under Table 7 — estimated cost to serve per successful hire with its major components — and one contribution-margin line in Table 8, labeled as an estimate the CRP will replace with measured data.

11. Reconcile institutional revenue with what Aim 3 actually produces. - *Section:* §8, Table 8; §6 post-CRP financing. - *Problem:* Year 4 books ~\$250K against a study design and explicitly non-causal modeled avoided cost. - *Reviewer inference:* The plan is selling evidence it has told the reviewer it will not have. - *Scoring consequence:* Commercialization Plan and Overall Impact. - *Minimum change:* Either reclassify Year 4 institutional revenue as a paid pilot or data/analytics contract priced on the evidence the CRP does produce (pathway analytics, county capacity data, time-to-care), or move first institutional revenue to Year 5 and state what the intervening proof-of-concept study costs and who funds it.

12. Add a provider-conversion funnel. - *Section:* §7, Table 6; §8. - *Problem:* 700 claimed listings and 3 lifetime payers, with no stated conversion assumption or required paying-account count per market. - *Reviewer inference:* The revenue model has a throughput input but no customer-count input. - *Minimum change:* State paying accounts required per market, the assumed claimed → paid conversion rate, and the assumed repeat-purchase rate per account per year.

13. Strengthen the IP section or reframe it honestly. - *Section:* §5. - *Problem:* Trade secret and copyright only; patents merely to be evaluated; the criterion explicitly asks about portfolio strength. - *Reviewer inference:* Nothing prevents replication once the approach is visible. - *Minimum change:* Name the two or three specific CRP-generated inventions that are patent candidates with a decision date, and state plainly that the primary moat is data and distribution rather than patents — an argued position beats an unstated one.

14. Make the Fundraising Plan concrete about the gap it must close. - *Section:* §6. - *Problem:* A \$3–5M raise is proposed against ~\$120K revenue by a company with no disclosed private financing history, on the strength of relationships twice described as non-commitments. - *Reviewer inference:* The financing plan is aspiration with a calendar attached. - *Minimum change:* State the specific metrics the applicant believes qualify it for that round (markets, run rate, retention, gross margin, institutional pipeline) and which CRP milestone produces each; and if the company has raised no private capital, say so and explain the strategic rationale rather than leaving reviewers to infer it from the 100% SBIR revenue table.

15. Restore accountability to the management milestones. - *Section:* §9, Table 9 footnote. - *Problem:* "Revenue values are planning projections rather than CRP success criteria" disclaims the only quantified milestones in the plan. - *Minimum change:* Keep the revenue figures as projections but add one committed, non-revenue threshold per year that is a criterion (e.g. hires per county per month, paying accounts per market, gross margin), so each gate has something that can fail.

Cross-document fixes required

- **Market count** — one number or one bounded rule, stated identically in both documents (ties to items 8 and 9).
- **Aim 1 verification method** — either implement blinded independent adjudication and describe it in Task 1.2, or remove "blinded expert review" from the Progress Report and "independently verifies" from Commercialization Plan §1.
- **Year 3 markets: same or new** — state explicitly whether the paid markets are the Year 2 free markets, new markets, or a stated split, and make Table 8 and Task 3.2 agree.
- **Pilot placements** — 20 or 25; use one number in both documents.
- **Price** — reconcile the \$250 base case with the observed \$50–\$275 per placement and \$275/month, and state whether the tested offer is per-hire, subscription, or both.
- **Independent financial validation** — make Research Strategy Table 3 and Commercialization Plan §9 agree; delete the "where commitment remains applicable" hedge or reflect it in both.
- **Component vocabulary** — use one tripartite decomposition in both documents, and give the Care Establishment Model a named presence in the Commercialization Plan.
- **Retention horizon** — remove 12-month retention from the CRP deliverables or state explicitly that only 3- and 6-month retention will be observable within the award.

Worth fixing if space permits

- Size the caregiver-recruitment market in dollars, not workers, in Commercialization Plan §4.
- Reconsider or reframe the 100/250/500-county illustration; as written it is arithmetically incompatible with the stated student wedge and invites the panel's most damaging calculation. If retained, state the additional labor pools and the yield assumptions that make it feasible.
- Address the free-to-paid transition risk in Year 3 (anchoring after a free Year 2) in one sentence.
- Add a decay/re-verification policy for the field-learning knowledge layer.

- Add one sentence on portal terms-of-use and account-risk handling for permissioned browser automation.
- Note the concurrent Phase IIB period (2021–2027) and how effort is separated across the overlap.
- State the referral-compliance position for per-placement fees paid by federally reimbursed providers; the plan already reasons about this for family referrals and the parallel is worth closing.
- Fix the Research Strategy's table numbering (two Table 1s, two Table 2s, jump to Table 6) and the "Tables 4 and 5" cross-reference.

Do not waste space fixing

- The absence of formal IP is not fixable in this cycle; argue the data-and-distribution moat well rather than manufacturing a patent story.
- Do not expand the competitive analysis. Research Strategy Table 2 and Commercialization Plan §4 are already better than the field standard.
- Do not add more citations to the unmet-need case. It is fully established, and additional evidence displaces the quantitative work that is actually missing.
- Do not restructure the Innovation section. The three-innovation framing works; only Innovation 3's evidentiary support is weak, and that is an Approach fix (test targeted against untargeted recruitment), not a rewrite.
- Do not soften the honest disclosures — the nil commercialization revenue, the non-commitment investor relationships, the "modeled rather than causal proof" language, or the willingness to publish a negative result. These are among the application's most credibility-generating passages, and reviewers explicitly credited them.
- Do not add detail to the Human Subjects procedures in the Research Strategy; those belong in the Human Subjects materials. Fix only the design-intrinsic items (severity stratification, adverse-event taxonomy, authorization model, independent monitoring).

Mock review prepared 2026-08-30. Scores reflect the Research Strategy and Commercialization Plan as supplied, under the stated assumption that absent application components are competent and reasonably supportive.